

A Model of Certifier and Accreditor Risk Calculation for Multi-Level Systems

Joe Loughry

Doctoral Student in the Department of Computer Science

University of Oxford

Wolfson Building, Parks Road

Oxford, OX1 3QD, UK

Email: joe.loughry@stx.ox.ac.uk

Tel: +1 303 221 4380

Abstract—From direct observation of the certification (post-software-development) and accreditation (pre-deployment) testing of cross domain systems used for the interconnection of classified security domains in U.S. and U.K. defence and intelligence community systems, certain characteristic behavioural patterns have been noted. The savvy developer can use these to exert a measure of control over the duration and cost of certification testing and to predict the likely direction and magnitude of the residual risk calculation performed by security accreditors in multi-lateral, multi-level, collateral, and compartmented site accreditations. DCID 6/3, Common Criteria, DIACAP, and ICD 503 testing efforts across the evolution of a long-lived software development programme were examined using grounded theory methodology. While discovered through investigation of classified cross domain system testing inefficiencies, it is believed that the principles are applicable more widely to privacy-sensitive areas such as electronic health care, financial, and law enforcement record keeping systems. The first thing found was a syndrome of pathological regressive interactions amongst software developers, managers, independent verification and validation contractors, penetration testers, and certification authorities that resulted in schedule slippage during the certification testing phase and, in the accreditation phase, ineffective duplication of testing with no corresponding improvement in residual risk. To understand why these problems occurred, an abstract model of how security accreditors discover and agree upon the true level of residual risk in multi-level cross domain system installations was developed. The model is powerful enough to handle collateral, SCI, and international cross domain systems with any number of endpoints. It works by establishing the visibility of threats, vulnerabilities, and mitigations from each data owner's perspective according to the associated accreditor's clearance over the space of all possible multi-level configurations, then identifying the smallest set of covert-channel-like information flows necessary to reach a concord about residual risk without violating the global security policy. Conventional wisdom holds that security rules should be strictly enforced, but it can be shown that under present regulations, some desirable information flows are inhibited and other undesirable information flows are forced. Paradoxically, it is sometimes the case that relaxing the rules actually improves security.

Index Terms—cross domain systems, certification and accreditation, security test and evaluation, certification test and evaluation.

I. INTRODUCTION

It is generally believed that certification and accreditation testing of cross domain systems takes more time than it ought

to and costs far more than necessary. Some of the reasons for this frustrating mis-allocation of resources lie in the observable duplication of effort that seems to be characteristic of cross domain solution security certification testing and cross domain system accreditation, as fully described in previous work [1], [2]. This paper describes progress toward a theory of cross domain security certification and accreditation deriving from long-term observation of a portion of the evolution of a representative cross domain solution for which software development began in 1992. Through more than twenty years of continuous maintenance and enhancement since, the software underwent numerous security evaluations by agencies with different testing criteria, aims, methods, and outcomes. Not all the outcomes were successful. It is believed that the contrast between successful and unsuccessful security testing activities of a succession of versions of the same software through multiple generations of security testing criteria offers a chance to find underlying principles that are more difficult to see in the typical instance of short-run product cycles that are otherwise characteristic of the cross domain software development ecosystem.

A. Certification and Accreditation of Cross Domain Systems

Cross domain systems are the backbone of interconnected networks of classified information processing systems. While not formally limited to classified information—the same principles apply to any multi-level security problem, for example the separation of engineering design information from marketing and accounting departments in a company, of financial trading and demand deposit functions in a merchant bank, chain of custody of evidence in law enforcement applications, or the accumulation, protection, and use of privacy-sensitive electronic health care records. This research, however, was done on the developers, installers, testers, and users of a cross domain solution designed for handling classified information in military and intelligence community environments.

Cross domain systems are what get installed in the field; they are built from components that include cross domain solutions. A cross domain solution is really a type of router, but often with additional capabilities at the Application layer atop the OSI model that go beyond 'deep packet inspection' to

include, sometimes, data format transliteration, content-aware message switching, and automatic sanitisation. Cross domain solutions are sometimes called ‘guards’.

1) *Certification*: Testing of cross domain solutions and cross domain systems is done in two phases independent of the developer’s own regression and functional tests. Every new and significantly updated versions of cross domain solution software must undergo certification testing and approval for use by the responsible government agency before it can be approved for use in handling classified information. Certification testing is usually performed by both Independent Verification and Validation (IV&V) contractors—who perform functional tests to verify that the product does everything it is supposed to do and nothing it is not supposed to do—and by penetration testers—commonly specialists employed by government agencies—who do more exotic and open-ended tests, searching for undiscovered vulnerabilities.

2) *Accreditation*: Once a cross domain solution has passed certification testing at the end of the development phase, it is approved for use in assembling cross domain systems, which are the things actually connected to customer networks. No cross domain system can be approved to operate and connect to classified information, however, before it has passed accreditation testing. Accreditation is the process of testing and approving a particular cross domain system, consisting of a particular cross domain solution connected to particular networks in a particular place, processing particular types of information, for a limited amount of time. Accreditation testing focusses on the details of particular data formats, transformations, rules, and communicating end-points. Because accreditation testing is done to the satisfaction of the accreditor—a government official whose job it is to accept personal responsibility for the correct operation of the cross domain system in production, and who is not the same person as the certification authority—accreditation testing is often seen to duplicate functional testing performed during the certification phase as well as, to some extent, specialised penetration testing that can be performed in the different and specific environment of the installation site.

B. Problems with Certification and Accreditation

It might be argued that the present level of security of these systems is adequate, and therefore that we know a reasonable upper bound on cost and time of security testing, but it may also be argued either that the level of security assurance is too high, impacting ship sailing schedules in the U.S. Navy; or too low, as evidenced by recent spills including the Bradley Manning and Edward Snowden cases, both of which *might* have been stopped by a cross domain guard [3, Ch. 1]. In addition, it is arguable that onerous certification and accreditation requirements discourage the introduction of innovative technologies to the problem, miring users in outdated, under-functioning systems for lack of alternatives.

What is known, however, is that certification and accreditation testing of cross domain solutions and systems is duplicative of effort at times; this is observable. What was not

understood until recently, however, is the interactions between people during the process that yield insight into their informal and formal communication channels. The present research aims to understand the reasons for security testing inefficiency with the goal of eliminating unnecessary redundancies in the process.

II. METHODOLOGY

The author was a participant observer in the software development and certification and accreditation testing of a leading cross domain solution product, from 1998–2012. During this time the software, here referred to as ‘*R*’, underwent certification testing at least seven times under three different criteria—DCID 6/3, Common Criteria (CC), and ICD 503—in addition to hundreds of accreditations. For three intervals in approximately 1999, 2004–6, and 2010–11, access to project records was obtained for research purposes. These encompass one abortive, one failed, and one completely successful security certification and accreditation effort of a single product.

Because the same group of people were largely engaged in developing and testing the software, through a series of versions, during all three of these times, it is believed that the methodology used here provides a measure of control over confounding factors. Controlled experiments are often unfeasible in software engineering due to the high cost of projects [4], [5]. For this reason, a grounded theory approach was used instead [6]. Grounded theory methodology is suitable to the form of the data available from the three projects, in which the author was a participant on the former projects and a non-participant observer on the latter; in exchange for access to telephone conference calls over the lifetime of the third project, the author minuted meetings between the developer and certifiers and accreditors and wrote some technical reports for the developer.

A. Participant Observation

In 1999 the author assisted with the preparation of ‘evaluable evidence’, of the type that would be used for a CC security evaluation, for the R_0 system. The intent was never to submit the evidence for validation or certification, but a Security Target (ST) was written describing the Security Functional Requirements to be claimed, and existing design documentation and software development process and life-cycle documentation collected in support of the Security Assurance Requirements specified at the Evaluation Assurance Level desired, comprising the evidence package.

Several years later, a foreign customer requested CC certification of the R' version of the software. The old evidence package was examined and found to be unusable because R' had new features that were not in R_0 , and besides lacked some functionality of R for export control reasons. The author began writing a new ST for R' and the latest process and life-cycle documentation was added to the evidence package. It was not so easy to locate up-to-date design documentation, however; the government sponsor of R had, in the intervening years, refused to pay for updating the design documents and

these had fallen out of date. Problems were encountered by the developer as a result during the validation phase. For reasons unrelated to R' , the customer unexpectedly cancelled the project shortly before the evaluation phase was scheduled to begin. CC certification was never achieved, although it should be emphasised that R did not fail security testing. The security evaluation effort, at a cost of approximately \$1.5 million, failed to achieve its objective.

B. Non-Participant Observation

The third set of project records was obtained for study in 2010–12 during the certification testing of R'' , a major version update to R . Particularly interesting about the effort this time was it was the first cross domain certification and accreditation done under a new set of testing criteria, ICD 503, which replaced DCID 6/3. None of the developer, certification authority, testers, or accreditors were familiar with the security controls (from NIST SP 800-53) and risk management framework (NIST SP 800-37) to be used. The author took minutes of numerous telephone conference calls involving the developer, certification authority, IV&V contractors, penetration testers, regression testers, and accreditors from the beginning of the R'' certification through the end of the first accreditation of a cross domain system incorporating R'' . The certification and accreditation of R'' were completely successful.

III. RESULTS

The grounded theory model derived from the data collected in Section II is shown in Figure 1 and presented fully in reference [3]. However, a brief overview of some of the more interesting results is possible here.

A strongly suggestive association was found between the occurrence of informal communications during certification testing activities and successful outcomes. Specifically, precursor communications carrying news of the results of regression testing and penetration testing in R'' often preceded planned formal communications, i.e., reports, by an interval sufficiently long to be helpful in immediate-term scheduling. Precursor communications were not seen in R' , although informal follow-on communications were. Furthermore, the R'' case study exhibited numerous examples of informal process improvement communication that were completely absent in R' .

Information asymmetry appeared to play a role in developer–certifier relations throughout the R'' certification, extending into the accreditation phase. Penetration testers, especially in a classified environment, as this was, are caught in a dilemma; they can inadvertently teach the developer to avoid poor practices at the cost of ‘using up’ some of the penetration testers’ proprietary vulnerability-finding techniques every time. This comes at a cost to the penetration testers, who must then replenish their store of new tricks through independent research. Obviously the goal of penetration testers in certification and accreditation testing is to harden the device under test, but this knowledge asymmetry between developer and tester evidently drives some of the personality conflicts

observed in the R'' data. Knowledge is continually transferred in multiple directions and not only through formal channels. In fact, belief appears to be qualitatively as influential as fact in the continuous updating of process knowledge and project knowledge that drive the decision making process controlling certification and accreditation.

IV. INTERPRETATION

It was just this sort of imbalance, as seen between the government security penetration testers and the developer during the certification phase that suggested the following approach applicable to the accreditation phase. What emerged was an abstract model of accreditor behaviour sufficient to describe all possible configurations of accreditors that have been seen in the field. Accreditors represent data owners and their associated systems that connect to a cross domain system. Accreditors perform a risk assessment decision based on their perceived value of ‘their’ information, their perceived effectiveness of the security controls implemented by the cross domain solution, and their perception of the threat environment.

This section describes the accreditor model and uses it to predict the behaviour of accreditors with various security clearances in their risk assessment decision that drives the willingness of each accreditor in a multi-level system to accept the residual risk of a particular accreditation.

A. Assumptions

In the model, accreditors have appropriate security clearances for their job. That is, for an accreditor who represents the interests of a data owner of SECRET information, that accreditor has a SECRET clearance and no higher. Similarly, an accreditor who represents the interests of a data owner with TOP SECRET information must have a TOP SECRET clearance. Under this assumption, which follows the Bell and LaPadula model strictly, an accreditor with a TOP SECRET clearance will never represent the interests of a data owner with information of a lower classification than TOP SECRET.

The assumption is necessary for the model for the model to work, but differs slightly from what happens in the real world. In practice, accrediting agencies tend to clear a pool of accreditors to the highest level applicable; this facilitates any accreditor to work on any accreditation, streamlining the personnel allocation problem of the government agency. It does, however, slightly degrade the quality of security protections because not all accreditations require TOP SECRET clearance and some accreditors have a higher security clearance than necessary.

The second assumption, for tractability, is that every cross domain system has exactly two accreditors. This assumption is not true in practice; real cross domain systems often affect more than two data owners, and in addition sometimes one accreditor represents the interests of two or more data owners. But to make the model visible through the notation of Venn diagrams, here it is assumed that there are two accreditors, arbitrarily designated the ‘high side’ and the ‘low side’,

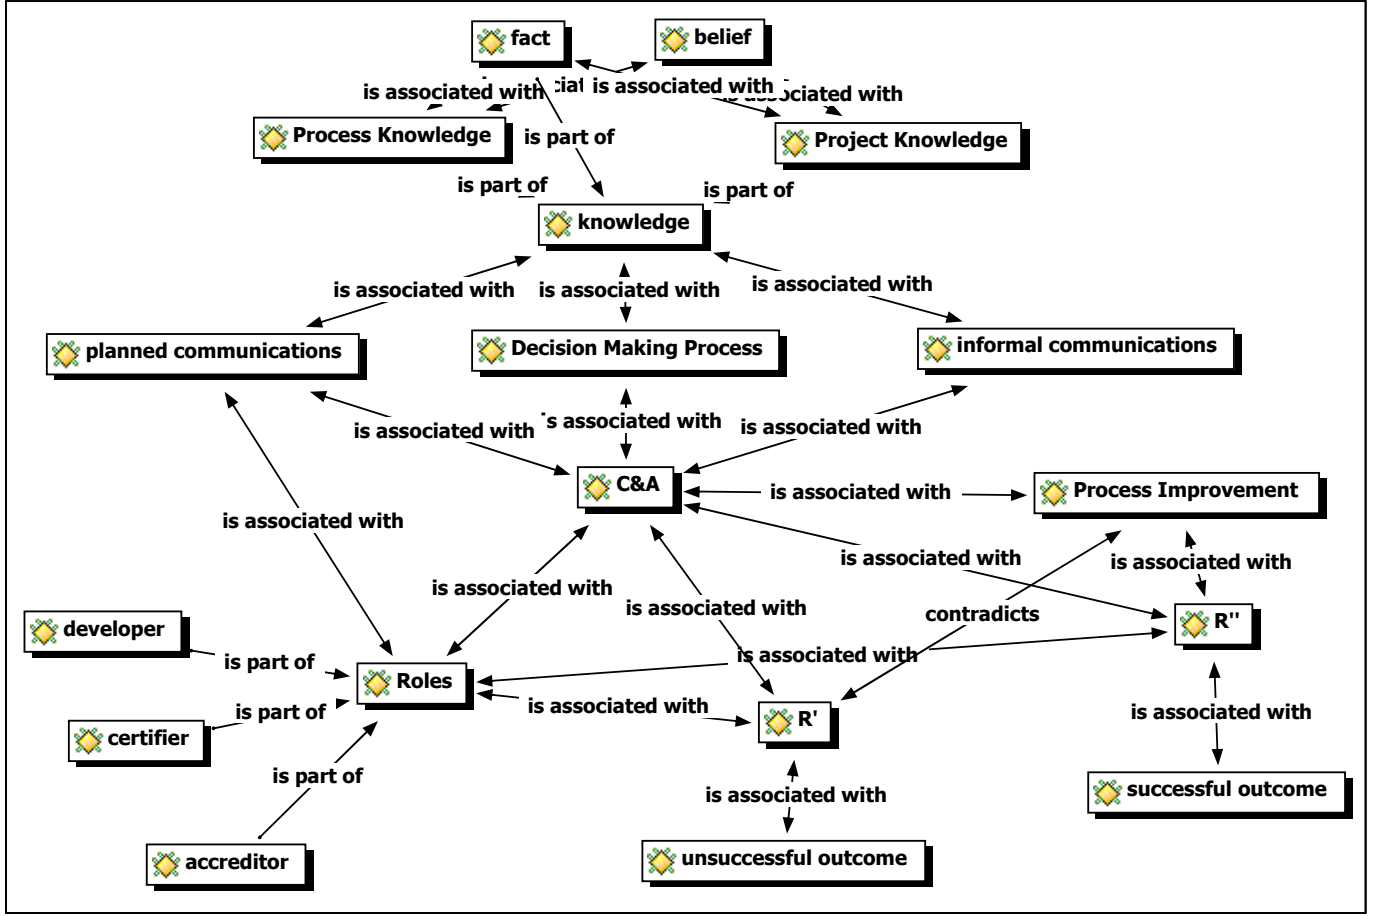


Figure 1: Grounded theory derived from data collected on the R_0 , R' , and R'' certification and accreditation projects.

although in Sensitive Compartmented Information (SCI)-like accreditations in general, and international ones particularly, the notion of which side of the cross domain system is the ‘high’ side is a matter of opinion. By definition, every cross domain system affects at least two data owners. To avoid the necessity of using Edwards’ spherical Venn diagrams, the number of accreditors is less than three [7]. Two is sufficient to represent every cross domain system accreditation, both collateral and SCI-like, including international accreditations, that the author has encountered in the field. For systems of more than two accreditors, they can be considered pair-wise without loss of generality.

B. Definitions

Security classification levels for classified information may be hierarchical or non-hierarchical. *Collateral* classifications are hierarchical, ranging from unclassified through CONFIDENTIAL, SECRET, and TOP SECRET, and are used by the military. A different type of classification, used by the intelligence community, is Sensitive Compartmented Information (SCI). SCI compartments are non-hierarchical and non-comparable, and the number of them is unlimited. TOP SECRET and SCI are often used in concert; for example, a

person with a TOP SECRET clearance would have access to CONFIDENTIAL information, but not SECRET CURVEBALL information, because CURVEBALL is (nominally) SCI. However, a person with SECRET FASTBALL CURVEBALL KNUCKLEBALL clearance would have access, assuming need-to-know.

C. Cross Domain System Configurations

Cross domain systems range in notional complexity from collateral, through mixed collateral-and-SCI, to pure SCI-like systems, and finally those containing one or more international accreditors. Figure 2 shows the relationship amongst the types; they can best be understood as a progression of overlapping knowledge. The larger circle, A_i , represents the set of security controls that Accreditor i thinks would be desirable to implement; this corresponds to the threat as perceived by i . The smaller circle, B_i , represents the set of security controls that Accreditor i believes possible or affordable to implement. The area between A_i and B_i is residual risk.

1) *Collateral Accreditations*: The simplest cross domain system is a collateral accreditation, meaning that all of the security classifications are related hierarchically. This is the the Bell-LaPadula model [8]. As shown in Figure 2(g), because

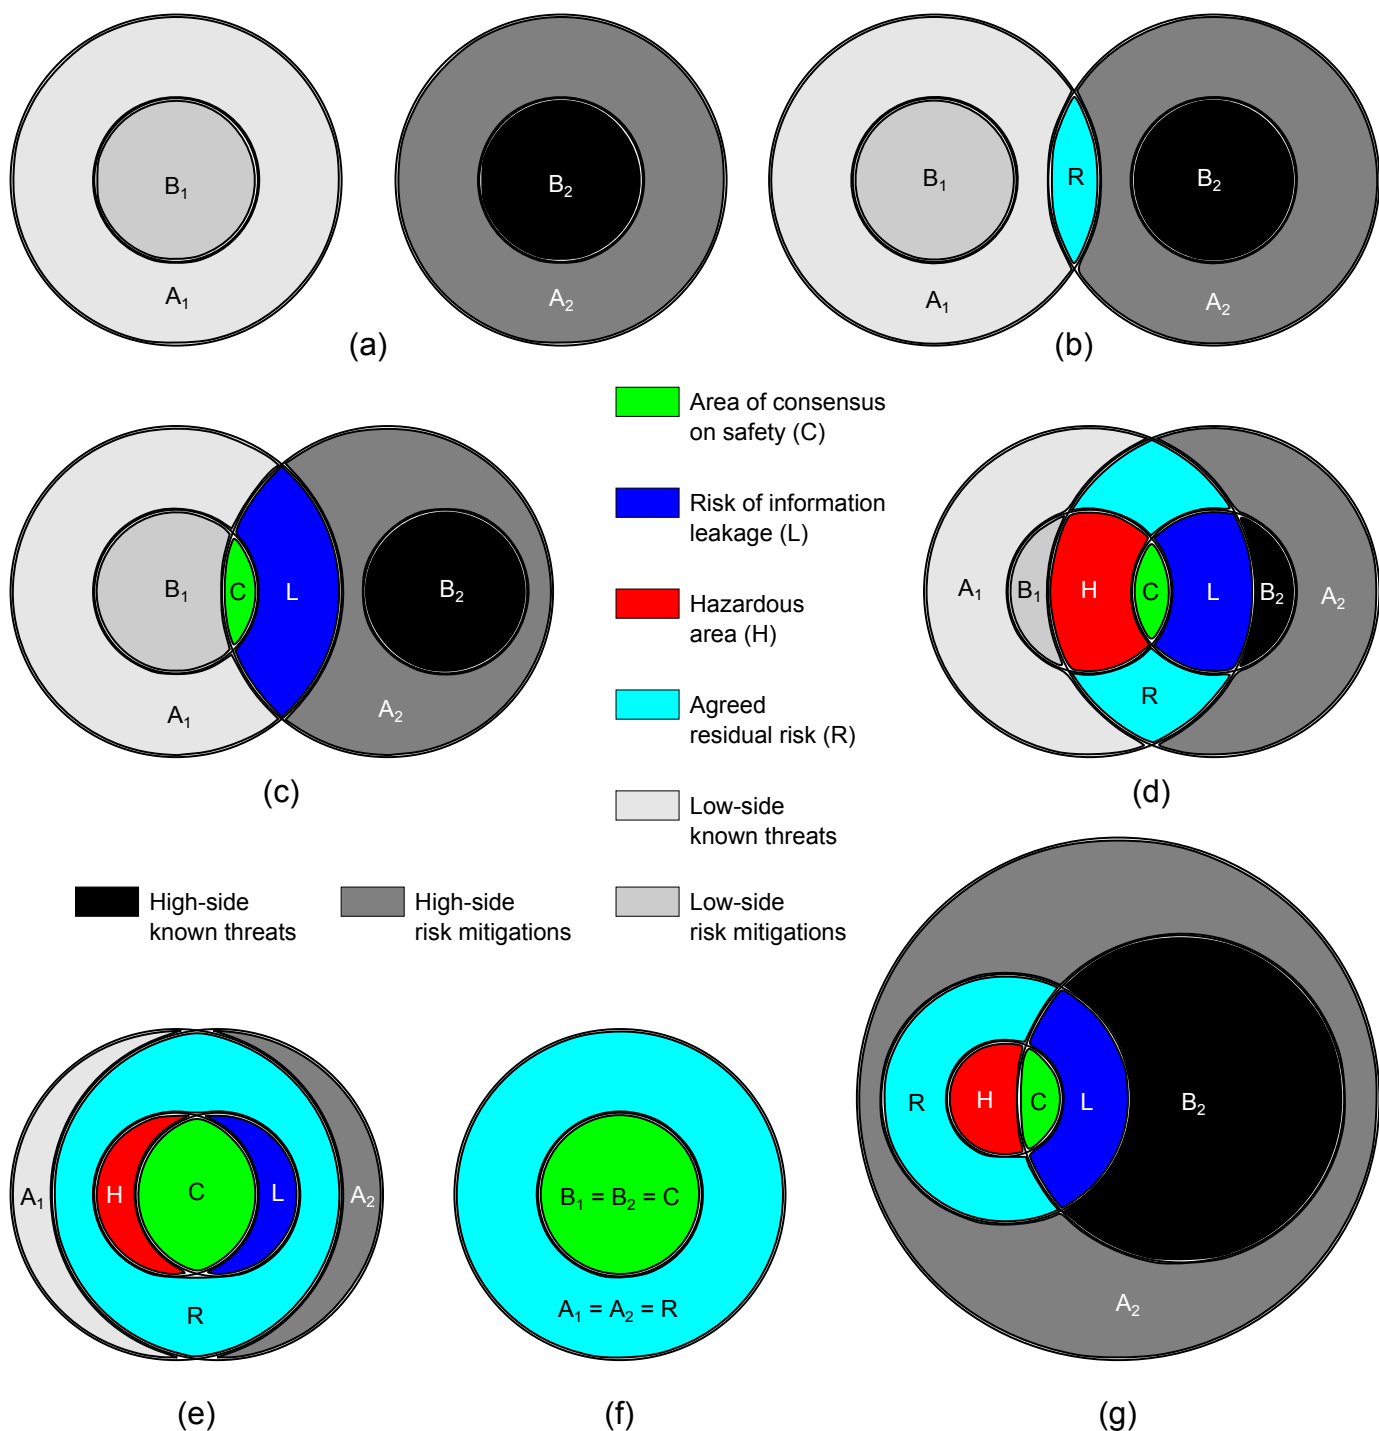


Figure 2: In sub-figure (a), which represents a pure SCI-type international accreditation, Accreditor 1 has a completely independent idea of what threats exist (A_1) and what risk mitigations are possible (A_2); similarly, Accreditor 2 is cognisant of a different set of risks (A_2) that would be desirable to mitigate, and a smaller set (B_2) of risks that are, from Accreditor 2's perspective, possible to mitigate. Sub-figure (b) represents public information. Sub-figure (c) represents classified information with a risk of information leakage. In sub-figure (d) there is not only a risk of information leakage but personal risk to Accreditor 2 as well. Sub-figure (e) shows how the risk of classified information leakage and personal risk to one or more accreditors diminishes as the situation approaches a pure collateral configuration. Sub-figure (f) is the degenerate configuration of a null cross domain system. Sub-figure (g) shows a collateral (non-SCI-like) accreditation in which Accreditor 1 has a lower security clearance than Accreditor 2; in this example, Accreditor 2 knows everything Accreditor 1 knows, but Accreditor 1 does not know everything Accreditor 2 knows.

both accreditors work for the same agency (e.g., the United States Department of Defence), it is often the case that the high-side accreditor may be aware of all of the lower-classified security threats and mitigations that the low-side accreditor is cleared to know about, but the high-side accreditor additionally may know of other, more highly classified threats or risk mitigations that the low-side accreditor is not cleared to know about. In the worst case, the high-side accreditor may even have the guilty knowledge that certain risk mitigations trusted by the low-side accreditor are ineffective in the face of highly classified threats. This leads to an information gradient from high to low that can, in some instances, force the high-side accreditor to reveal classified information to the low-side accreditor (coloured red in Figure 2(g)) in violation of the high-side accreditor's oath to protect classified information. In other instances, coloured blue in Figure 2(g), the low-side accreditor may be able to deduce certain information above the low-side accreditor's security clearance from the fact that the high-side accreditor refuses to accept the residual risk of the cross domain system accreditation without saying why; this leaks information.

2) *SCI Accreditations*: In an uncomplicated SCI accreditation with no shared compartments, shown in Figure 2(a), there is no risk of information leakage; each accreditor performs the risk assessment calculation independently. In Figure 2(b) there is public information; both accreditors agree that region C is an acceptable risk. In Figure 2(c), there is a risk of information leakage from Accreditor 1 to Accreditor 2, because Accreditor 1 considers region C to pose an unacceptable risk, and by insisting on the security controls in region C , inadvertently communicates to Accreditor 2 the existence of a classified threat mitigation that Accreditor 2 is not cleared to know about. Figure 2(d) is where it gets interesting. In configuration (d), both accreditors may be aware of SCI threat mitigations that they are not allowed to tell the other accreditor about, but in addition at least one accreditor may have guilty knowledge that some threat mitigations (in region H) believed effective by the other accreditor do not work.¹

As the overlap grows between SCI accreditors, as in Figure 2(e), the hazardous region H shrinks, until the degenerate case of Figure 2(f) where both sides of the cross domain system have the same security clearance and classification.

3) *SCI-like and International Accreditations*: In the real world, things do not always seem so clear-cut. Cross domain systems interconnecting international partners, or unclassified systems (such as news feeds) are not uncommon. It can readily be seen, however, that the accreditor model accommodates such complications with ease, as both international accreditors and unclassified endpoints can simply be treated as SCI systems that happen to share no compartments. The model

¹Alternatively, this may be interpreted as the knowledge of both the existence of a classified threat and the fact that there is no known mitigation for it. When this situation occurs in a collateral accreditation, there is no risk of information leakage because from the perspective of the high-side accreditor, it is a risk he or she cannot talk about, and for the low-side accreditor, ignorance is bliss.

is therefore complete.

V. SUMMARY AND CONCLUSION

Cross domain systems, by definition, inhabit a niche where they necessarily always involve at least two mutually distrustful data owners. Consequently, the security certification and accreditation testing of components and systems will always tend to be rigorous. Previously it has been shown that apparently unnecessary duplication of effort exists in the testing of cross domain solutions and cross domain systems, but now it is beginning to be understood why this has occurred. A grounded theory analysis of several variably successful case studies of cross domain certification and accreditation projects led to the development of an abstract model of accreditor behaviour suggesting that some disagreements, by reason of individual perspective from different sides of a cross domain system, must lead to some amount of duplicated testing effort in order to reach a consensus between accreditors with different—and sometimes incompatible—security clearances.

VI. ACKNOWLEDGEMENTS

Grateful acknowledgement is hereby given to the Lockheed Martin Corporation for access to project records and data. The author wishes particularly to thank his supervisors, Dr Andrew Martin and Dr Ivan Fléchais for their insightful criticism of the accreditor model and its applicability.

VII. AUTHOR BIOGRAPHY

Joe Loughry received the B.A. degree in mathematics from the University of Colorado, Boulder in 1986 and the M.S.E. degree in software engineering from Seattle University in 1996. He is currently a Ph.D. student in computer science at the University of Oxford. He did some of the first research on compromising optical emanations and holds U.S. patent 6,987,461 on countermeasures. His research interests include cross domain systems, security certification and accreditation testing, penetration testing, side channels, and efficient enumeration of subsets.

REFERENCES

- [1] J. Loughry, "Unsteady ground: Certification to unstable criteria," in *Proceedings of the Second International Conference on Advances in System Testing and Validation Lifecycle*, Nice, France, 22–27 August 2010.
- [2] —, "Information asymmetry in classified cross domain system security accreditation," in *Computer and Electronics Security Applications Rendez-vous (C&ESAR) 2012*, Rennes, France, 20–22 November 2012, pp. 19–28.
- [3] —, "Security test and evaluation of cross domain systems," Ph.D. dissertation, University of Oxford, Michaelmas Term 2013.
- [4] G. M. McCue, "IBM's Santa Teresa Laboratory—architectural design for program development," *IBM Systems Journal*, vol. 17, no. 1, pp. 4–25, 1978.
- [5] J. Rost and R. L. Glass, *The Dark Side of Software Engineering: Evil on Computing Projects*. Los Alamitos, California: IEEE Computer Society Press, 2011.
- [6] B. G. Glaser and A. L. Strauss, *The Discovery of Grounded Theory: Strategies for Qualitative Research*. Aldine Transaction, 1967.
- [7] A. W. Edwards, *Cogwheels of the Mind: The Story of Venn Diagrams*. Baltimore, Maryland: Johns Hopkins University Press, 2004.
- [8] D. E. Bell and L. J. LaPadula, "Secure computer systems: Mathematical foundations," The MITRE Corporation, Technical Report 2547, Volume I, March 1973.